

Московский государственный технический университет им. Н.Э. Баумана
Кафедра «Системы обработки информации и управления»



Лабораторная работа №4
по дисциплине
«Методы машинного обучения»
на тему

«Реализация алгоритма Policy Iteration»

Выполнил:
студент группы ИУ5И-22М
Лю Чжинань

Москва — 2023 г.

1. Цель лабораторной работы

ознакомление с базовыми методами обучения с подкреплением.

2. Задание

На основе рассмотренного на лекции примера реализуйте алгоритм Policy Iteration для любой среды обучения с подкреплением (кроме рассмотренной на лекции).

3. Текст программы

```
from matplotlib import pyplot as plt

import numpy as np

from scipy.stats import poisson

import seaborn as sns

"""
Hyper parameters

<int> MAX_NUM: 每个停车场能停的最大数目 (20)

<int> lam1_rent: 停车场 1 租车 $\lambda$ 值 (3)

<int> lam1_return: 停车场 1 还车 $\lambda$ 值 (3)

<int> lam2_rent: 停车场 2 租车 $\lambda$ 值 (4)

<int> lam2_return: 停车场 2 还车 $\lambda$ 值 (2)

<int> MAX_ACTION: 最大移动汽车数目 (5)

<int> CAR_COST: 移动车辆的代价 (2)

<int> CAR_EARNING: 租车的收入 (10)

<float> DISCOUNT: 收益折扣 (0.9)

<np.array> actions: 动作集合 (-5, -4, ..., 4, 5)

<int> POISSON_UPPER_BOUND: 限制泊松分布产生请求数目的上限
```

```

<dict> poisson_cache: 存储每个 (n,  $\lambda$ ) 对应的泊松概率, key 为 n*(POISSON_UPPER_BOUND-1)+lam
"""

MAX_NUM = 20

lam1_rent = 3

lam1_return = 3

lam2_rent = 4

lam2_return = 2

MAX_ACTION = 5

DISCOUNT = 0.9

CAR_COST = 2

CAR_EARNING = 10

actions = np.arange(-MAX_ACTION, MAX_ACTION + 1)

POISSON_UPPER_BOUND = 11

poisson_cache = dict()

```

```

def poisson_prob(n, lam):

    global poisson_cache

    key = n * (POISSON_UPPER_BOUND - 1) + lam

    if key not in poisson_cache:

        poisson_cache[key] = poisson.pmf(n, lam)

    return poisson_cache[key]

```

```

class dp:

    def __init__(self):

        self.v = np.ones((MAX_NUM + 1, MAX_NUM + 1), float)

        self.actions = np.zeros((MAX_NUM + 1, MAX_NUM + 1), int)

```

```
self.gama = DISCOUNT

self.delta = 0

self.theta = 0.01

pass
```

```
def state_value(self, state, action, state_value, constant_returned_cars):

    """
```

```
:param state: 状态定义为每个地点的车辆数

:param action: 车辆的移动数量[-5,5]，负：2->1，正：1->2

:param state_value: 状态价值矩阵

:param constant_returned_cars: 将换车的数目设定为泊松均值，替换为泊松概率分布

:return:

    """

# initial total return

returns = 0.0
```

```
# 移动车辆产生负收益

returns -= CAR_COST * abs(action)
```

```
# 移动后的车辆总数不能超过 20

NUM_OF_CARS_1 = min(state[0] - action, MAX_NUM)

NUM_OF_CARS_2 = min(state[1] + action, MAX_NUM)
```

```
# 遍历两地全部的可能概率下（截断泊松概率）租车请求数目

for rent_1 in range(POISSON_UPPER_BOUND):
```

```

        for rent_2 in range(POISSON_UPPER_BOUND):

            # prob 为两地租车请求的联合概率，概率为泊松分布
            prob = poisson_prob(rent_1, lam1_rent) * poisson_prob(rent_2,
lam2_rent)

            # 两地原本汽车数量
            num_of_cars_1 = NUM_OF_CARS_1
            num_of_cars_2 = NUM_OF_CARS_2

            # 有效租车数目必须小于等于该地原有的车辆数目
            valid_rent_1 = min(num_of_cars_1, rent_1)
            valid_rent_2 = min(num_of_cars_2, rent_2)

            # 计算回报，更新两地车辆数目变动
            reward = (valid_rent_1 + valid_rent_2) * CAR_EARNING

            num_of_cars_1 -= valid_rent_1
            num_of_cars_2 -= valid_rent_2

            # 如果还车数目为泊松分布的均值
            if constant_returned_cars:

                # 两地的还车数目均为泊松分布均值
                returned_cars_1 = lam1_return
                returned_cars_2 = lam2_return

                # 还车后总数不能超过车场容量
                num_of_cars_first_loc = min(num_of_cars_1 + returned_cars_1,
MAX_NUM)

                num_of_cars_second_loc = min(num_of_cars_2 + returned_cars_2,
MAX_NUM)

                # 核心：
                # 策略评估:  $V(s) = p(s', r | s, \pi(s)) [r + \gamma V(s')]$ 
                returns += prob * (reward + DISCOUNT *
state_value[num_of_cars_first_loc, num_of_cars_second_loc])

```

```

        # 否则计算所有泊松概率分布下的还车空间

        else:

            for returned_cars_first_loc in range(POISSON_UPPER_BOUND):

                for returned_cars_second_loc in
range(POISSON_UPPER_BOUND):

                    probb_return = poisson_prob(

                        returned_cars_first_loc, lam1_return) *

poisson_prob(

                        returned_cars_second_loc, lam2_return)

                    num_of_cars_first_loc_ = min(num_of_cars_1 +
returned_cars_first_loc, MAX_NUM)

                    num_of_cars_second_loc_ = min(num_of_cars_2 +
returned_cars_second_loc, MAX_NUM)

                    # 联合概率为【还车概率】*【租车概率】

                    probb_ = probb_return * probb

                    returns += probb_ * (reward + DISCOUNT *

                                                                    state_value[num_of_cars_first_loc
_, num_of_cars_second_loc_])

            return returns

```

```

def policy_iteration(self, constant_returned_cars=True):

```

```

    """

```

```

        :param constant_returned_cars:

```

```

        :return:

```

```

    """

```

```

    # 设置迭代参数

```

```

    iterations = 0

```

```

# 准备画布大小，并准备多个子图

_, axes = plt.subplots(2, 3, figsize=(40, 20))

# 调整子图的间距，wspace=0.1 为水平间距，hspace=0.2 为垂直间距

plt.subplots_adjust(wspace=0.1, hspace=0.2)

# 这里将子图形成一个 1*6 的列表

axes = axes.flatten()

while True:

    # 使用 seaborn 的 heatmap 作图

    # flipud 为将矩阵进行垂直角度的上下翻转，第 n 行变为第一行，第一行变为第 n 行，
    如此。

    # cmap:matplotlib 的 colormap 名称或颜色对象；

    # 如果没有提供，默认为 cubehelix map（数据集为连续数据集时）或 RdBu_r（数
    据集为离散数据集时）

    fig = sns.heatmap(np.flipud(self.actions), cmap="rainbow",
    ax=axes[iterations])

```

```

# 定义标签与标题

fig.set_ylabel('# cars at first location', fontsize=30)

fig.set_yticks(list(reversed(range(MAX_NUM + 1))))

fig.set_xlabel('# cars at second location', fontsize=30)

fig.set_title('policy {}'.format(iterations), fontsize=30)

```

```

# policy evaluation (in-place) 策略评估 (in-place)

# 未改进前，第一轮 policy 全为 0，即[0, 0, 0...]

while True:

    old_value = self.v.copy()

    for i in range(MAX_NUM + 1):

```

```

        for j in range(MAX_NUM + 1):
            # 更新  $V(s)$ 
            new_state_value = self.state_value([i, j], self.actions[i, j], self.v, constant_returned_cars)

            # in-place 操作
            self.v[i, j] = new_state_value

            # 比较  $V_{old}(s)$ 、 $V(s)$ ，收敛后退出循环
            max_value_change = abs(old_value - self.v).max()

            print('max value change {}'.format(max_value_change))

            if max_value_change < 1e-4:
                break

```

```

# policy improvement

# 在上一部分可以看到，策略 policy 全都是 0，如不进行策略改进，其必然不会收敛到实际最优策略。

# 所以需要如下策略改进

policy_stable = True

# i、j 分别为两地现有车辆总数

for i in range(MAX_NUM + 1):
    for j in range(MAX_NUM + 1):
        old_action = self.actions[i, j]

        action_returns = []

        # actions 为全部的动作空间，即 [-5、-4...4、5]

        for action in actions:
            if (0 <= action <= i) or (-j <= action <= 0):
                action_returns.append(self.state_value([i, j], action, self.v, constant_returned_cars))
            else:

```



```

        action_returns.append(-np.inf)

        # 找出产生最大动作价值的动作
        new_action = actions[np.argmax(action_returns)]

        # 更新策略
        self.actions[i, j] = new_action

        if policy_stable and old_action != new_action:
            policy_stable = False

    print('policy stable {}'.format(policy_stable))

```

```

    if policy_stable:
        fig = sns.heatmap(np.flipud(self.v), cmap="rainbow", ax=axes[-1])
        fig.set_ylabel('# cars at first location', fontsize=30)
        fig.set_yticks(list(reversed(range(MAX_NUM + 1))))
        fig.set_xlabel('# cars at second location', fontsize=30)
        fig.set_title('optimal value', fontsize=30)

        break

```

```

        iterations += 1

        # plt.title('Policy Iteration')

        plt.savefig('./policy_iteration.png')

        plt.show()

        plt.close()

    return

```

```

def value_iteration(self, constant_returned_cars=True):
    """

```

```
:param constant_returned_cars:

:return:

"""

# 设置迭代参数

iterations = 0
```

```
# 准备画布大小，并准备多个子图

_, axes = plt.subplots(1, 3, figsize=(40, 10))

# 调整子图的间距，wspace=0.1 为水平间距，hspace=0.2 为垂直间距

plt.subplots_adjust(wspace=0.1, hspace=0.2)

# 这里将子图形成一个 1*6 的列表

axes = axes.flatten()

while True:

    # 使用 seaborn 的 heatmap 作图

    # flipud 为将矩阵进行垂直角度的上下翻转，第 n 行变为第一行，第一行变为第 n 行，
    如此。

    # cmap:matplotlib 的 colormap 名称或颜色对象；

    # 如果没有提供，默认为 cubehelix map（数据集为连续数据集时）或 RdBu_r（数
    据集为离散数据集时）

    fig = sns.heatmap(np.flipud(self.actions), cmap="rainbow",
ax=axes[iterations])
```

```
# 定义标签与标题

fig.set_ylabel('# cars at first location', fontsize=30)

fig.set_yticks(list(reversed(range(MAX_NUM + 1))))

fig.set_xlabel('# cars at second location', fontsize=30)

fig.set_title('policy {}'.format(iterations), fontsize=30)
```

```
# value iteration 价值迭代
```

```
while True:
```

```
    old_value = self.v.copy()

    for i in range(MAX_NUM + 1):
        for j in range(MAX_NUM + 1):
            action_returns = []

            # actions 为全部的动作空间，即[-5、-4...4、5]
            for action in actions:
                if (0 <= action <= i) or (-j <= action <= 0):
                    action_returns.append(self.state_value([i, j],
action, self.v, constant_returned_cars))
                else:
                    action_returns.append(-np.inf)

            # 找出产生最大动作价值的动作
            max_action = actions[np.argmax(action_returns)]

            # 更新 V(s)
            new_state_value = self.state_value([i, j], max_action,
self.v, constant_returned_cars)

            # in-place 操作
            self.v[i, j] = new_state_value

            # 比较 V_old(s)、V(s)，收敛后退出循环
            max_value_change = abs(old_value - self.v).max()

            print('max value change {}'.format(max_value_change))

            if max_value_change < 1e-4:
                print(iterations)

                break
```

```

# policy improvement

policy_stable = True

# i、j 分别为两地现有车辆总数

for i in range(MAX_NUM + 1):
    for j in range(MAX_NUM + 1):
        old_action = self.actions[i, j]

        action_returns = []

        # actions 为全部的动作空间，即[-5、-4...4、5]

        for action in actions:
            if (0 <= action <= i) or (-j <= action <= 0):
                action_returns.append(self.state_value([i, j], action,
self.v, constant_returned_cars))
            else:
                action_returns.append(-np.inf)

        # 找出产生最大动作价值的动作

        new_action = actions[np.argmax(action_returns)]

        # 更新策略

        self.actions[i, j] = new_action

        if policy_stable and old_action != new_action:
            policy_stable = False

print('policy stable {}'.format(policy_stable))

```

```

if policy_stable:
    fig = sns.heatmap(np.flipud(self.v), cmap="rainbow", ax=axes[-1])
    fig.set_ylabel('# cars at first location', fontsize=30)
    fig.set_yticks(list(reversed(range(MAX_NUM + 1))))

```

```
fig.set_xlabel('# cars at second location', fontsize=30)

fig.set_title('optimal value', fontsize=30)

break
```

```
iterations += 1
```

```
# plt.title('Value Iteration')

plt.savefig('./value_iteration.png')

plt.show()

plt.close()

return
```

```
if __name__ == '__main__':

    model = dp()

    model.policy_iteration(constant_returned_cars=True)
```

```
model1 = dp()

model1.value_iteration(constant_returned_cars=True)

pass
```

4. Результат

